

Import Libraries

```
In [76]: # Import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

%matplotlib inline
sns.set()
```

```
In [77]: conda install -c https://conda.anaconda.org/plotly plotly
```

Collecting package metadata (current_repodata.json): done
Solving environment: done

All requested packages already installed.

Note: you may need to restart the kernel to use updated packages.

load dataset

```
In [78]: # reading data with pandas
df = pd.read_csv('Car_details.csv')
```

```
In [79]: # detect how many rows and columns
df.shape
```

```
Out[79]: (8127, 13)
```

```
In [80]: df.head()
```

```
Out[80]:
```

	name	year	selling_price	km_driven	fuel	seller_type	transi
0	Skoda Rapid 1.5 TDI Ambition	2014	370000	120000	Diesel	Individual	
1	Honda City 2017- 2020 EXi	2006	158000	140000	Petrol	Individual	
2	Hyundai i20 Sportz Diesel	2010	225000	127000	Diesel	Individual	
3	Maruti Swift VXI BSIII	2007	130000	120000	Petrol	Individual	
4	Hyundai Xcent 1.2 VTVT E Plus	2017	440000	45000	Petrol	Individual	

Explore data analysis

```
In [81]:
```

```
#here we can see the Dtype of each column and check if th
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8127 entries, 0 to 8126
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   name                   8127 non-null   object
1   year                   8127 non-null   int64
2   selling_price          8127 non-null   int64
3   km_driven              8127 non-null   int64
4   fuel                   8127 non-null   object
5   seller_type            8127 non-null   object
6   transmission           8127 non-null   object
7   owner                  8127 non-null   object
8   mileage                7906 non-null   object
9   engine                 7906 non-null   object
10  max_power              7912 non-null   object
11  torque                 7905 non-null   object
12  seats                  7906 non-null   float64
dtypes: float64(1), int64(3), object(9)
memory usage: 825.5+ KB

```

```

In [82]: #displays summary statistics
         df.describe()

```

```

Out[82]:

```

	year	selling_price	km_driven	seats
count	8127.000000	8.127000e+03	8.127000e+03	7906.000000
mean	2013.803987	6.382950e+05	6.981020e+04	5.416772
std	4.044497	8.063003e+05	5.654780e+04	0.959637
min	1983.000000	2.999900e+04	1.000000e+00	2.000000
25%	2011.000000	2.549990e+05	3.500000e+04	5.000000
50%	2015.000000	4.500000e+05	6.000000e+04	5.000000
75%	2017.000000	6.750000e+05	9.800000e+04	5.000000
max	2020.000000	1.000000e+07	2.360457e+06	14.000000

```

In [83]: df.columns

```

```

Out[83]: Index(['name', 'year', 'selling_price', 'km_driven', 'fuel',
                'seller_type',
                'transmission', 'owner', 'mileage', 'engine', 'max_power',
                'torque',
                'seats'],
                dtype='object')

```

```
In [84]: #show the data types for each columns  
df.dtypes
```

```
Out[84]: name          object  
year            int64  
selling_price    int64  
km_driven        int64  
fuel            object  
seller_type      object  
transmission     object  
owner           object  
mileage          object  
engine           object  
max_power        object  
torque           object  
seats            float64  
dtype: object
```

```
In [85]: # check how many null values in our dataset  
df.isna()
```

```
Out[85]:
```

	name	year	selling_price	km_driven	fuel	seller_type	transr
0	False	False	False	False	False	False	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	
...	
8122	False	False	False	False	False	False	
8123	False	False	False	False	False	False	
8124	False	False	False	False	False	False	
8125	False	False	False	False	False	False	
8126	False	False	False	False	False	False	

8127 rows × 13 columns

```
In [86]: #here we can see that the last 5 cloumn has null values  
df.isna().sum()
```

```
Out[86]: name          0
        year          0
        selling_price  0
        km_driven     0
        fuel          0
        seller_type    0
        transmission   0
        owner         0
        mileage       221
        engine        221
        max_power     215
        torque        222
        seats         221
        dtype: int64
```

```
In [87]: #we can see the percentage null values from the dataset
        df.isna().sum()/raw_data.shape[0]
```

```
Out[87]: name          0.000000
        year          0.000000
        selling_price  0.000000
        km_driven     0.000000
        fuel          0.000000
        seller_type    0.000000
        transmission   0.000000
        owner         0.000000
        mileage       0.027193
        engine        0.027193
        max_power     0.026455
        torque        0.027316
        seats         0.027193
        dtype: float64
```

Cleaning

```
In [88]: #we gonna fill the seats column with mean(float64)
        df.seats.fillna(df.seats.mean(),inplace=True)
```

```
In [89]: df.seats.mean()
```

```
Out[89]: 5.416772071844173
```

```
In [90]: #we gonna fill the engine column with mode(object)
        df.engine.fillna(df.engine.mode()[0],inplace=True)
```

```
In [91]: df.engine.mode()
```

```
Out[91]: 0      1248 CC  
dtype: object
```

```
In [92]: #we gonna fill the engine column with mode(object)  
df.mileage.fillna(df.mileage.mode()[0],inplace=True)#we g
```

```
In [93]: df.mileage.mode()
```

```
Out[93]: 0      18.9 kmpl  
dtype: object
```

```
In [94]: # drop the 'torque', 'max_power'  
df.drop(columns=['torque', 'max_power'],inplace=True)
```

```
In [95]: #now we can see there is no more empty valuse  
df.isna().sum()
```

```
Out[95]: name          0  
year          0  
selling_price  0  
km_driven     0  
fuel          0  
seller_type   0  
transmission  0  
owner         0  
mileage       0  
engine        0  
seats         0  
dtype: int64
```

```
In [96]: #cheack if there is duplicated and we found 1202  
df.duplicated().sum()
```

```
Out[96]: 1202
```

```
In [97]: #drop the duplicated  
df.drop_duplicates(inplace=True)
```

```
In [98]: #notice that after we clean our dataset the num of the co  
df.shape
```

```
Out[98]: (6925, 11)
```

visualization

```
In [99]: #count of selling prices in every year  
# plotting the bar chart  
fig = px.bar(df,x="year", y='selling_price')  
  
# showing the plot  
fig.show()
```

In [100...

```
#count of selling prices according to it's km_driven  
# plotting the histogram  
fig = px.histogram(df, x="km_driven", y="selling_price")  
  
# showing the plot  
fig.show()
```


In [101...

```
#sum of selling_price for each cars  
# plotting the bar chart  
fig = px.histogram(df,x="name", y='selling_price')  
  
# showing the plot  
fig.show()
```

In [102...

```
#count of selling prices in every year using countplot  
plt.figure(figsize=[15,8])  
sns.countplot(df.year,hue='selling_price',data=df)
```

/opt/anaconda3/lib/python3.9/site-packages/seaborn/_decorators.py:36: FutureWarning:

Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

Out[102...


<AxesSubplot:xlabel='year', ylabel='count'>





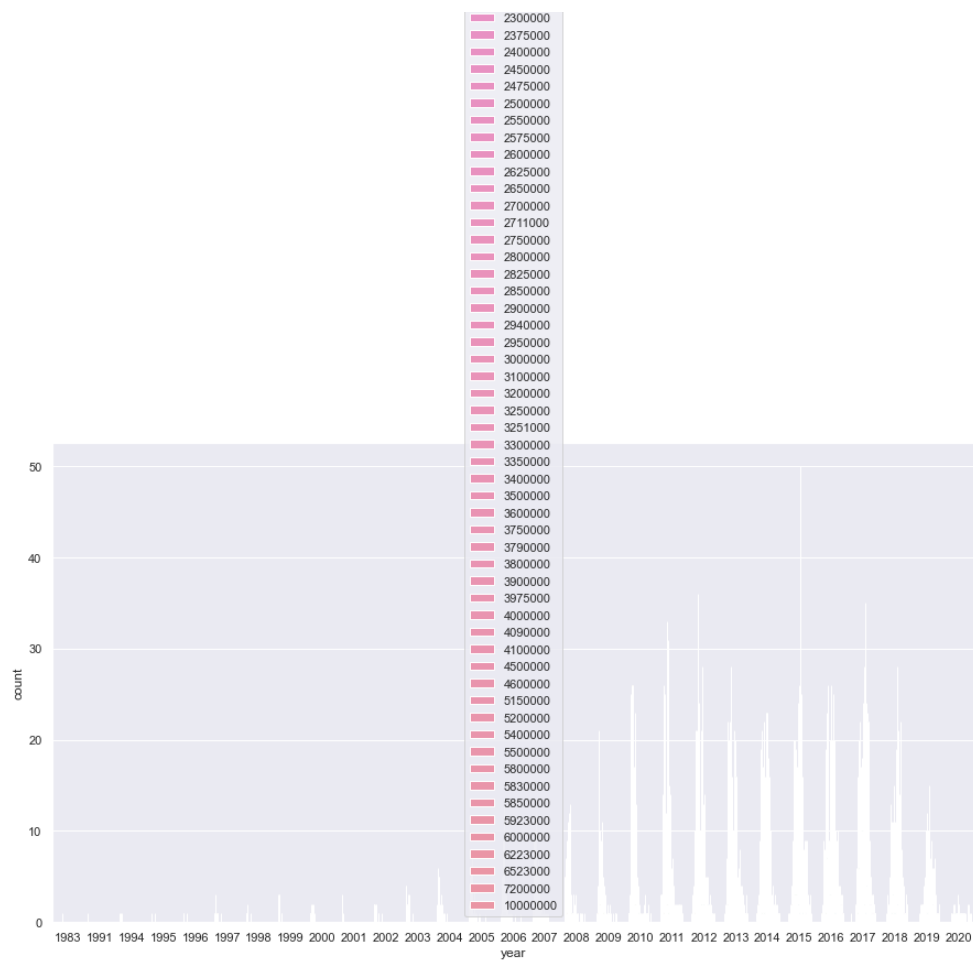
	179000
	180000
	181000
	184000
	185000
	187000
	189000
	190000
	191000
	194000
	195000
	196000
	198000
	199000
	200000
	202999
	204000
	204999
	206000
	209000
	210000
	211000
	212000
	215000
	217000
	219000
	220000
	221000
	222000
	225000
	227000
	227999
	229999
	231000
	235000
	236000
	238000
	240000
	241000
	242000
	244000
	245000
	246000
	248000
	250000
	250999
	252000
	252999
	254000
	254999
	256000
	256999
	258000
	259000
	260000
	261000
	262000
	265000
	267000
	268000
	270000
	272000
	275000
	276000
	277000
	278000
	279000
	280000
	282000
	285000
	287000
	290000
	291000
	292000
	295000
	297000
	298000
	299000
	300000
	302000
	305000
	307000
	308000
	310000
	311000
	312000
	313000
	315000
	316000
	317000
	318000
	319000
	320000
	321000
	325000
	327000
	328000
	329000
	330000
	335000
	337000
	339000
	340000
	344000





725000
730000
731000
735000
736000
737000
740000
741000
745000
746000
749000
750000
751000
752000
754000
755000
756000
757000
760000
764000
765000
770000
773000
775000
778000
779000
780000
785000
786000
790000
791000
793000
795000
797000
799000
800000
801000
802999
803999
805000
808000
810000
811000
813000
819999
822000
825000
830000
833000
834000
835000
839000
840000
844999
849000
850000
851000
860000
861000
861999
866000
869999
875000
880000
885000
889000
890000
891000
892000
894999
899000
900000
905000
906000
910000
911000
911999
919999
925000
930000
934000
940000
944999
946000
949000
950000
955000
957000
960000
965000
969999
974000
975000
977000
978999
980000
990000
999000
1000000
1019999
1025000
1030000
1031000
1035000

1040000
1044999
1050000
1051000
1075000
1080000
1085000
1090000
1100000
1110000
1125000
1132000
1140000
1143000
1145000
1147000
1149000
1150000
1151000
1160000
1165000
1175000
1180000
1190000
1200000
1210000
1220000
1225000
1227000
1237000
1250000
1251000
1265000
1270000
1275000
1282000
1290000
1295000
1300000
1325000
1350000
1365000
1380000
1385000
1390000
1400000
1405000
1425000
1445000
1450000
1460000
1465000
1475000
1490000
1500000
1511000
1515000
1516000
1520000
1525000
1530000
1550000
1560000
1575000
1576000
1590000
1594000
1600000
1625000
1630000
1650000
1670000
1675000
1680000
1689999
1700000
1745000
1748999
1750000
1757000
1789999
1800000
1825000
1850000
1859000
1864999
1888000
1898999
1900000
1920000
1925000
1938000
1950000
2000000
2051000
2064000
2100000
2125000
2150000
2175000
2199000
2200000
2280000




In [103...

```
# plotting the countplot for the yaer
plt.figure(figsize=[15,8])
sns.countplot(df.year)
```

```
/opt/anaconda3/lib/python3.9/site-packages/seaborn/_decorators.py:36: FutureWarning:
```

Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

Out[103.. <AxesSubplot:xlabel='year', ylabel='count'>



In [104...

```
#find the correlation between variables  
df.corr()
```

Out[104...

	year	selling_price	km_driven	seats
year	1.000000	0.433080	-0.377070	0.023098
selling_price	0.433080	1.000000	-0.165615	0.157154
km_driven	-0.377070	-0.165615	1.000000	0.205931
seats	0.023098	0.157154	0.205931	1.000000

In [105...

```
# visualize the correlations between the variables  
sns.heatmap(df.corr(), cmap="seismic", annot=True, vmin=-
```



```
In [106... #to visualize the relationship between each variable
plt.figure(figsize=[20,30])
sns.pairplot(df,diag_kind='kde',hue='selling_price',height=
```

```
Out[106... <seaborn.axisgrid.PairGrid at 0x7ff8823005b0>
```



preprocessing for modeling

```
In [107... #convert character columns to dummy variables to be ready
df=pd.get_dummies(df)
```

```
In [108... df.head()
```

```
Out[108...
      year  selling_price  km_driven  seats  name_Ambassador  name_A
      year  selling_price  km_driven  seats  CLASSIC 1500 DSL  Classic
      AC
0  2014      370000      120000      5.0      0
1  2006      158000      140000      5.0      0
2  2010      225000      127000      5.0      0
3  2007      130000      120000      5.0      0
4  2017      440000      45000      5.0      0
```

5 rows × 2590 columns

```
In [109... #you can see that the columns increase after dummies vari
df.shape
```

```
Out[109... (6925, 2590)
```

split Data

```
In [110... # detect input and output
x=df.drop('selling_price',axis=1)
y=df.selling_price
y
```

```
Out[110...
0      370000
1      158000
2      225000
3      130000
4      440000
...
8120    260000
8121    475000
8122    320000
8123    135000
8124    382000
Name: selling_price, Length: 6925, dtype: int64
```

```
In [111... # train test split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(x, y,
```

```
In [112... X_train
```

```
Out[112...
          year km_driven seats name_Ambassador name_Ambassador
          CLASSIC 1500 DSL Classic 2000 DSZ
                        AC AC PS
2464 2008      58632    5.0              0              0
2437 2011      50000    5.0              0              0
7515 2018      50000    5.0              0              0
3063 2006      62900    5.0              0              0
1236 2013     300000    8.0              0              0
...   ...         ...    ...              ...              ...
5701 1997     100000    5.0              0              0
3620 2008     110000    5.0              0              0
1817 2013      70000    9.0              0              0
2844 2017      90000    7.0              0              0
2981 2016      59300    5.0              0              0
```

5540 rows x 2589 columns

```
In [113... y_train
```

```
Out[113... 2464      114999
2437      250000
7515     5500000
3063      135000
1236      200000
...
5701       70000
3620      250000
1817      500000
2844     1050000
2981      240000
Name: selling_price, Length: 5540, dtype: int64
```

In [114...

```
x_test
```

Out[114...

	year	km_driven	seats	name_Ambassador CLASSIC 1500 DSL AC	name_Ambassador Classic 2000 DSZ AC PS
1173	2012	70000	5.0	0	0
942	2008	100000	5.0	0	0
4282	2016	94000	5.0	0	0
313	2014	80100	5.0	0	0
5397	2011	25000	5.0	0	0
...
4104	2008	136500	8.0	0	0
351	2012	40000	5.0	0	0
2575	2019	14000	5.0	0	0
3350	2017	44000	5.0	0	0
4274	2016	120000	5.0	0	0

1385 rows × 2589 columns

In [115...

```
y_test
```

Out[115...

```
1173    254999
942      290000
4282    1100000
313      555000
5397     245000
...
4104     375000
351      325000
2575     380000
3350    2600000
4274     350000
Name: selling_price, Length: 1385, dtype: int64
```

```
In [116... # feature scaling
#scale all the futur to same scale
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

Modeling

```
In [117... # apply LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

```
In [118... linreg_model = LinearRegression()
linreg_model.fit(x,y)
```

```
Out[118... LinearRegression()
```

```
In [119... y_pred = linreg_model.predict(x)
print(metrics.mean_absolute_error(y, y_pred))
```

```
44834.248419519965
```

```
In [120... print(linreg_model.intercept_)
print(linreg_model.coef_)
```

```
3405192073097.6074
[ 3.07951709e+04 -2.07833909e-01 -9.01974974e+04 ... -3.7
 7938069e+11
 -4.43547773e+11 -1.81815062e+11]
```

Model evaluation

```
In [121... # R-Squared
linreg_model.score(x, y)
```

```
Out[121... 0.9741883150866637
```

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js